

M4L3 Homework Assignment

Joshua Conte

October 15, 2017

1 M4L3 Homework Assignment

R studio was configured with the following parameters before beginning the project:

```
# clears the console in RStudio
cat("\014")
```

```
# clears environment
rm(list = ls())

# Load required packages
require(ggplot2)
require(cluster)
require(ama)
require(useful)
require(mclust)
```

1.1 Load Data.

I opened the Wholesale customers Data Set using read.csv2 and downloaded it directly from the UC Irvine Machine Learning Repository.

To format the data, the data is separated by ',', stringsAsFactors = FALSE so that the strings in a data frame will be treated as plain strings and not as factor variables. I set na.strings for missing data. Once the data was loaded I added the column names and changed the data types to numeric and finally removed the text data type.

Below is my R code:

```
# Some csv files are really big and take a while to open. This command checks to
# see if it is already opened, if it is, it does not open it again.
# I also omitted the first column
if (!exists("dfWCD")) {
dfWCD <-
  read.csv2("Wholesale customers data.csv",
    sep = ",",
    stringsAsFactors = FALSE,
    na.strings=c("", "NA")
  )
# Add a column so I know which study the data is refereing to
study <- sprintf("study_%s",seq(1:440))
dfWCD$study<-study
}

# Download directly from site (unreliable from Ecuador)
# if (!exists("dfWCD")) {
# dfWCD <-
#   read.csv2(
#     url(
#       "https://archive.ics.uci.edu/ml/machine-learning-databases/00292/Wholesale customers data.csv"
#     ),
#     sep = ",",
#     stringsAsFactors = FALSE,
#     na.strings=c("", "NA")
#   )
```

```

# # Add a column so I know which study the data is refereing to
# study <- sprintf("study_%s",seq(1:440))
# dfWCD$study<-study
# }

# change 2 to 24 to numeric
dfWCD[1:8] <- sapply(dfWCD[1:8], as.numeric)

# Print first lines
str(dfWCD)

## 'data.frame': 440 obs. of 9 variables:
## $ Channel : num 2 2 2 1 2 2 2 1 2 ...
## $ Region : num 3 3 3 3 3 3 3 3 3 ...
## $ Fresh : num 12669 7057 6353 13265 22615 ...
## $ Milk : num 9656 9810 8808 1196 5410 ...
## $ Grocery : num 7561 9568 7684 4221 7198 ...
## $ Frozen : num 214 1762 2405 6404 3915 ...
## $ Detergents_Paper: num 2674 3293 3516 507 1777 ...
## $ Delicassen : num 1338 1776 7844 1788 5185 ...
## $ study : chr "study_1" "study_2" "study_3" "study_4" ...

# Select the first 8 lines for plotting
dfWCD2<-dfWCD[1:8]

# Print first lines
str(dfWCD2)

## 'data.frame': 440 obs. of 8 variables:
## $ Channel : num 2 2 2 1 2 2 2 1 2 ...
## $ Region : num 3 3 3 3 3 3 3 3 3 ...
## $ Fresh : num 12669 7057 6353 13265 22615 ...
## $ Milk : num 9656 9810 8808 1196 5410 ...
## $ Grocery : num 7561 9568 7684 4221 7198 ...
## $ Frozen : num 214 1762 2405 6404 3915 ...
## $ Detergents_Paper: num 2674 3293 3516 507 1777 ...
## $ Delicassen : num 1338 1776 7844 1788 5185 ...

```

1.2 Expectation-maximization (EM)

The expectation-maximization (EM) algorithm is an iterative method for finding maximum likelihood or maximum a posteriori (MAP) estimates of parameters in statistical models, where the model depends on unobserved latent variables. Expectation Maximization (EM) is perhaps most often used algorithm for unsupervised learning.

EM clustering probabilistically assigns data to different clusters. This is sometimes called “soft-clustering” (as opposed to “hard-clustering” in which data only belongs to one cluster).

R has the `Mclust` function from the `mclust` library to provide the estimating the parameters in a statistical model using the EM algorithm.

1.2.1 Mclust Process

Given a model (a mixture Guassians, Binomial, etc.), we have the following parameters:

- X : This is a set of observed data. (Doesn't change)
- Z : This is a set of estimates for unobserved values
- T : This is a set of unknown parameters for our model

The expectation-maximization steps:

- Initialize the unknown parameters T to random values.
- Compute the best fit for the missing values Z using the existing parameter values.
- Use the best fit missing values Z to generate a better estimate for the unknown parameters T

Iterate until we have a convergence, usually when Z and T don't change much or after a fixed number of steps.

1.2.2 Mclust Example with Wholesale customers Data Set

Note that the summary command with an mclust object generates:

- log.likelihood: This is the log likelihood of the BIC value
- n: This is the number of X points
- df: This is the degrees of freedom
- BIC: This is the Bayesian information criteria; low is good
- ICL: Integrated Complete X Likelihood-a classification version of the BIC.

This data has 2 outputs, channel and region. Below, I am going to use region as my output to analyze and remove both channel and region and cluster the remaining data.

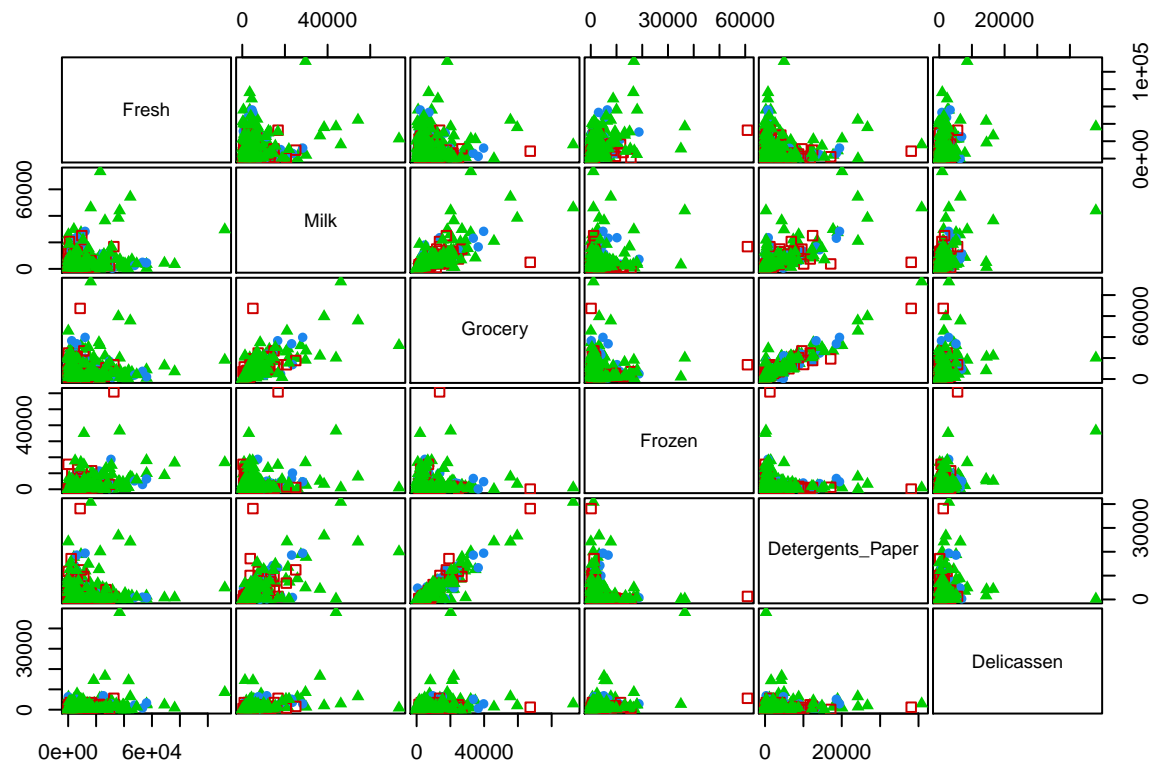
```
# I am going to use region
region.d = dfWCD2$Region
table(region.d)
```

```
## region.d
##    1    2    3
## 77  47 316
```

```
# I am going to remove region and channel from the data since they are not needed.
X = dfWCD2[3:8]
head(X)
```

```
##    Fresh Milk Grocery Frozen Detergents_Paper Delicassen
## 1 12669 9656    7561    214                2674    1338
## 2  7057 9810    9568   1762                3293    1776
## 3  6353 8808    7684   2405                3516    7844
## 4 13265 1196    4221   6404                 507    1788
## 5 22615 5410    7198   3915                1777    5185
## 6  9413 8259    5126    666                1795    1451
```

```
# Cluster Plot
clPairs(X, region.d)
```



```
# This fits the data
fit <- Mclust(X)
fit

## 'Mclust' model object:
## best model: diagonal, varying volume and shape (VVI) with 8 components

summary(fit)


## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
##
## Mclust VVI (diagonal, varying volume and shape) model with 8 components:
##
## log.likelihood  n  df      BIC      ICL
##      -24319.1 440 103 -49265.15 -49349.14
##
## Clustering table:
##   1  2  3  4  5  6  7  8
## 66 17 55 87 14 36 124 41
```

The plot command for EM produces the following four plots:

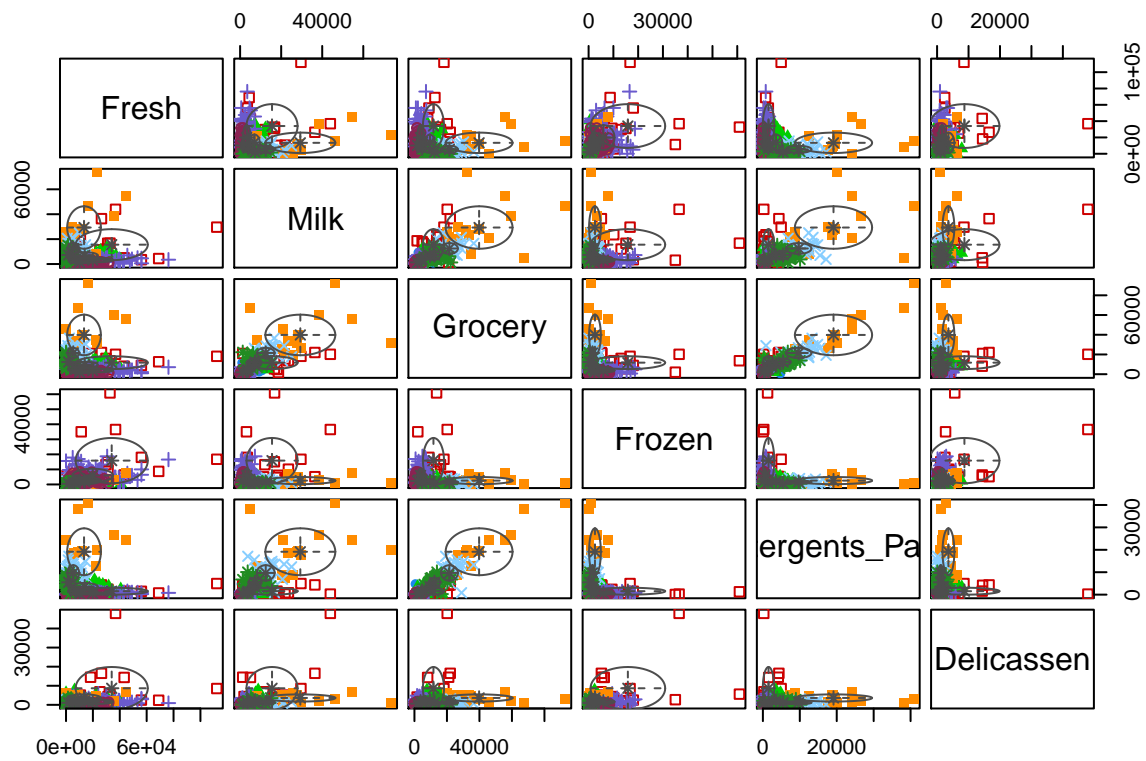
- The BIC values used for choosing the number of clusters
- A plot of the clustering

- A plot of the classification uncertainty
- The orbital plot of clusters

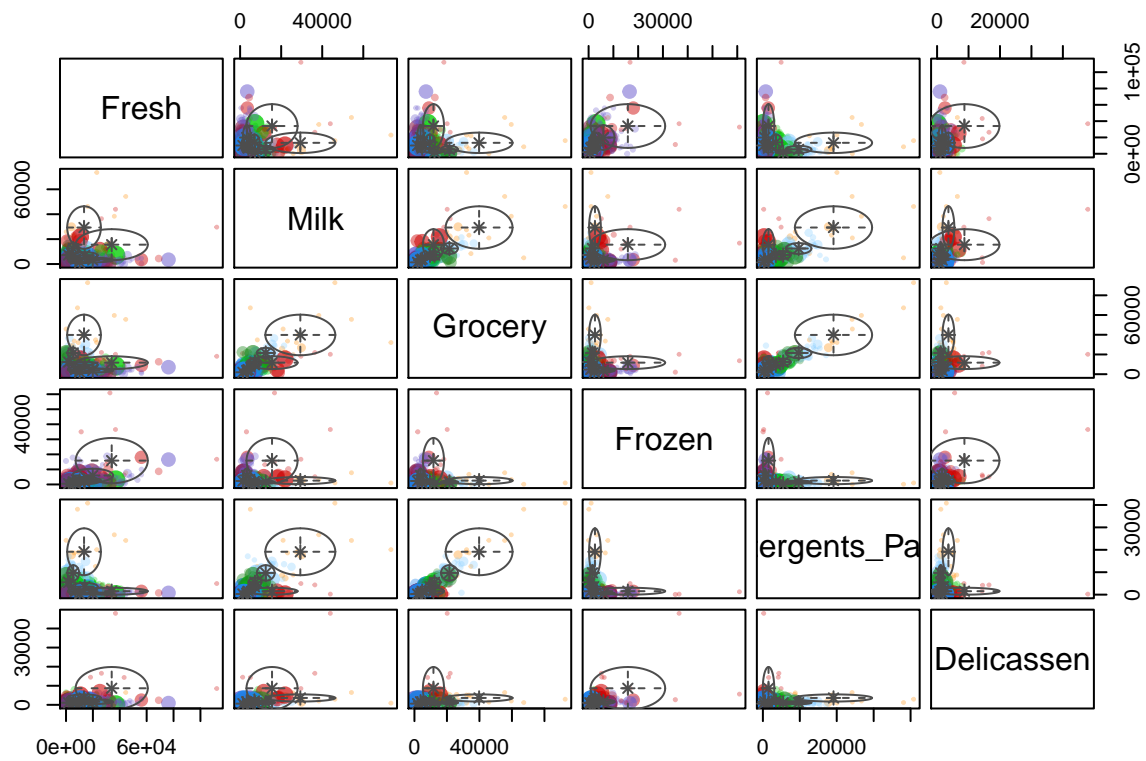
```
# 1: BIC
plot(fit, what = "BIC")
```



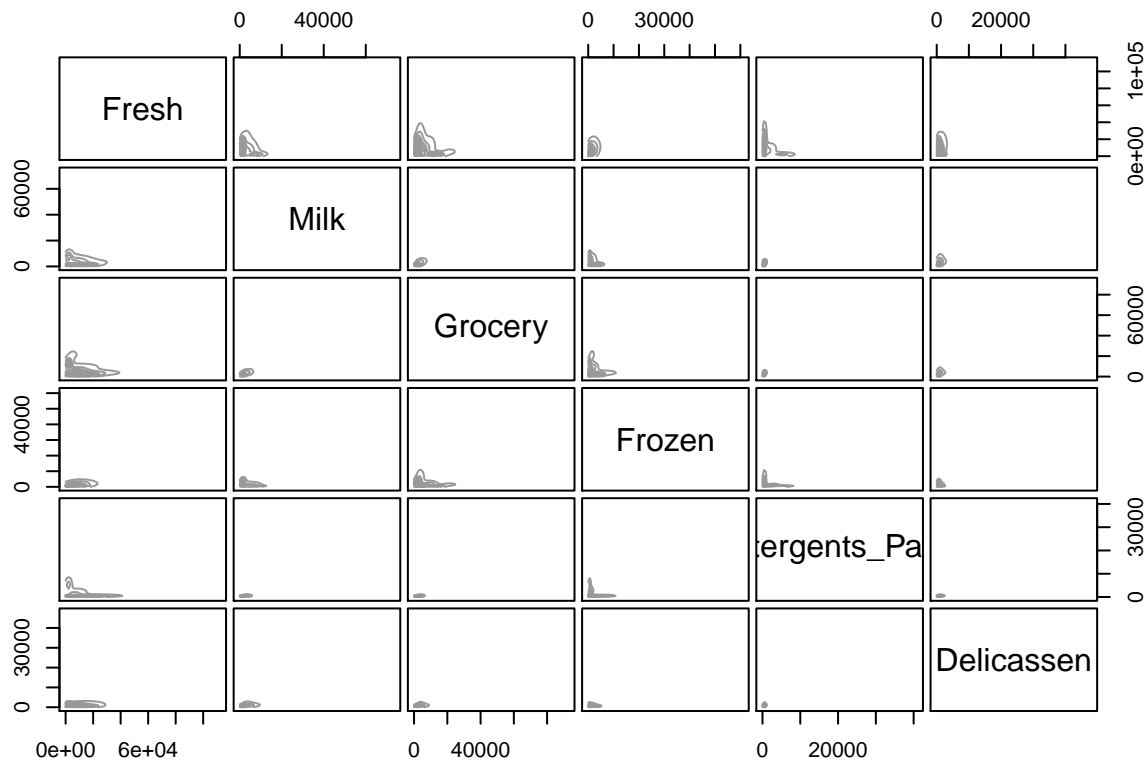
```
#table(class, fit$BIC)
#2: classification
plot(fit, what = "classification")
```



```
#table(class, fit$classification)
# 3: uncertainty
plot(fit, what = "uncertainty")
```



```
#table(class, fit$uncertainty)
# 4: density
plot(fit, what = "density")
```


```
#table(class, fit$density)
```

The `mclustBIC` from `mclust` uses BIC for EM initialized by model-based hierarchical clustering for parameterized Gaussian mixture models.

```
BIC = mclustBIC(X)
summary(BIC)
```

```
## Best BIC values:
##           VVI,8      VVI,7      VVI,9
## BIC      -49265.15 -49317.38677 -49365.1714
## BIC diff       0.00    -52.24061   -100.0252
```

```
plot(BIC) # Only BIC plot
```



```
# plot(BIC,what = "BIC")
```


```
# The following just get BIC plot: plot(BIC,what = "classification"), plot(BIC,what = "uncertainty"), p
```

mclustICL from mclust uses ICL (Integrated Complete-data Likelihood) for parameterized Gaussian mixture models fitted by EM algorithm initialized by model-based hierarchical clustering.

```
ICL = mclustICL(X)
summary(ICL)
```

```
## Best ICL values:
##           VVI,8      VVI,7      VVI,6
## ICL      -49349.14 -49399.15142 -49447.53960
## ICL diff      0.00   -50.01335   -98.40153
```

```
plot(ICL) # Only ICL plot
```



##M4L2 Analysis ###K-means This calculation makes the k-means calculation more stable, it performs this analysis 1000 times and takes the ones with the least error:

```
# Clustering with 4 clusters
```

```
k <- 8
```

```
trails<-1000
```

```
X.8.cluster <- kmeans(X,k, nstart = trails)
```

```
X.8.cluster
```

```
## K-means clustering with 8 clusters of sizes 2, 3, 32, 88, 5, 174, 30, 106
```

```
##
```

```
## Cluster means:
```

```
##      Fresh      Milk  Grocery  Frozen Detergents_Paper Delicassen
## 1 34782.000 30367.000 16898.000 48701.500      755.500 26776.0000
## 2 85779.667 12503.667 12619.667 13991.667      2159.000 3958.0000
## 3 38912.094 4767.750 5452.344 5081.000      811.625 2129.5312
## 4 4398.580 8712.114 12551.375 1445.636      5331.125 1507.0568
## 5 25603.000 43460.600 61472.200 2636.000     29974.200 2708.8000
## 6 5708.356 2463.247 3061.029 2699.379      794.454 846.1092
## 7 6683.067 17468.033 26658.933 1986.300     11872.900 2531.2000
## 8 18860.425 3423.858 4793.991 3584.934      1136.462 1585.3585
```

```
##
```

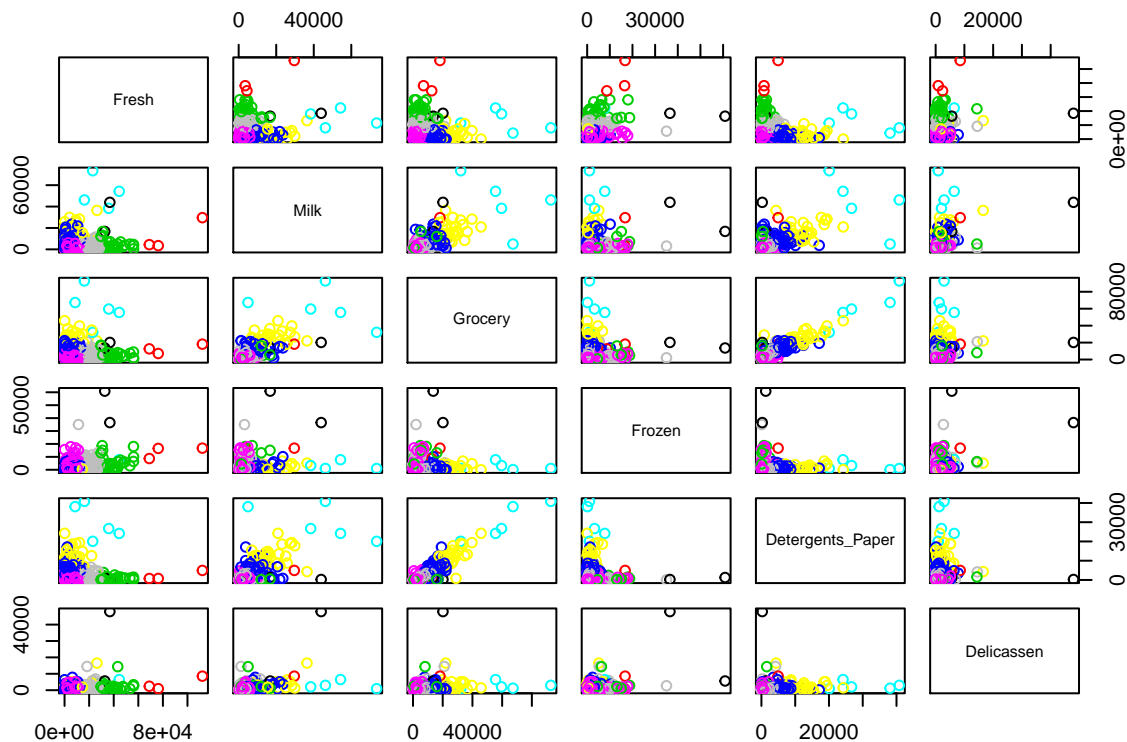
```
## Clustering vector:
```

```
## [1] 8 4 4 8 8 6 8 4 6 4 4 8 3 8 8 6 4 6 8 6 8 6 3 7 8 8 6 8 7 3 8 6 8 3 6
## [36] 4 3 4 4 3 8 8 4 7 4 7 7 5 4 7 6 6 3 4 8 6 7 4 8 4 6 5 6 4 6 7 6 8 6 6
## [71] 8 8 6 8 4 8 6 7 6 6 6 4 4 8 6 5 5 3 6 8 6 8 7 8 4 6 6 6 6 6 4 4 4 3 8
## [106] 8 4 4 4 7 6 4 8 8 8 6 6 6 8 6 8 6 6 4 3 2 8 8 6 3 6 6 8 6 6 6 4 4 8 6
## [141] 8 3 3 6 8 7 6 6 6 3 8 6 8 6 6 4 4 8 4 4 4 6 8 7 4 7 4 6 6 6 4 7 6 4 6
```


```
## [176] 4 3 8 6 6 8 2 4 1 6 6 6 4 4 4 8 8 6 4 6 8 3 4 6 6 7 7 8 6 6 7 6 6 6 4
## [211] 8 7 6 4 4 4 7 8 4 6 8 4 6 6 6 8 8 6 6 6 6 8 6 8 6 6 8 6 3 8 8 8 6 4
## [246] 4 6 8 8 6 6 7 6 3 4 3 6 6 3 3 6 6 8 6 4 4 4 8 4 8 6 6 4 3 6 6 8 6 6 8
## [281] 6 8 3 8 2 3 6 8 8 3 6 6 6 4 8 6 8 6 4 6 8 7 4 4 7 4 7 8 6 4 6 3 4 6 6
## [316] 4 6 6 6 7 6 6 8 8 8 1 6 6 8 6 6 7 8 5 8 8 8 6 6 6 4 4 6 7 6 6 4 8 6 7
## [351] 6 7 6 4 8 6 8 4 4 6 8 6 6 6 6 4 6 6 8 6 3 8 6 8 6 6 4 3 6 4 8 8 3 6 4
## [386] 6 6 8 6 6 6 6 6 8 6 6 4 6 6 6 6 8 8 8 8 6 8 4 6 6 6 6 4 6 6 4 4 4 6
## [421] 4 8 8 8 8 6 4 3 6 6 4 6 8 6 8 3 3 7 6 6
##
## Within cluster sum of squares by cluster:
## [1] 1591649631 1657529737 4300328413 5126690922 5682449098 5813885754
## [7] 5004238144 6732045153
## (between_SS / total_SS = 77.2 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"
## [5] "tot.withinss" "betweenss"    "size"         "iter"
## [9] "ifault"
```

These are the plots:

```
plot(X,col=X.8.cluster$cluster)      # Plot Clusters
```



```
# Centroid Plot against 1st two discriminant functions
clusplot(X, X.8.cluster$cluster, color=TRUE, shade=TRUE, labels=2, lines=0)
```



These two components explain 72.46 % of the point variability.

```
# library(fpc)
# plotcluster(dfWCD2, dfWCD2.4.cluster$cluster)
```

1.2.3 PAM

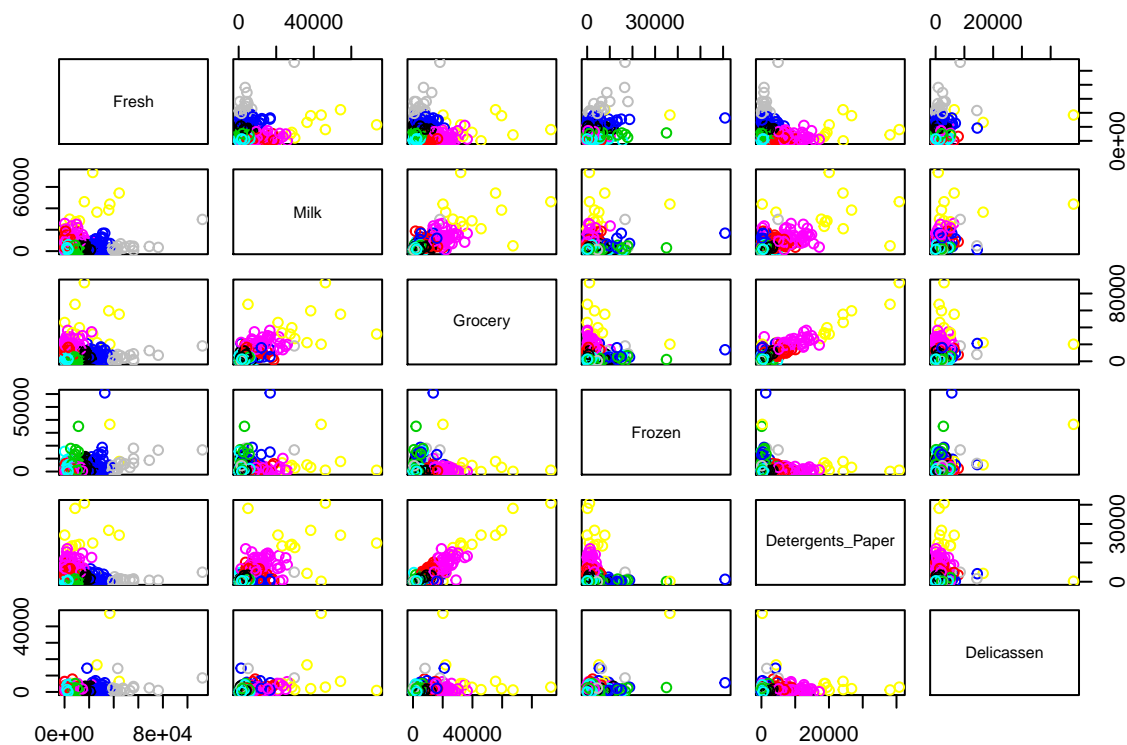
PAM Analysis:

```
# PAM
k<-8
X.pam.8.clust<- pam(X,k, keep.diss = TRUE, keep.data = TRUE)
X.pam.8.clust
```

```
## Medoids:
##      ID Fresh  Milk Grocery Frozen Detergents_Paper Delicassen
## [1,] 121 17160 1200   3412   2417           174         1136
## [2,] 198  2427  7097  10391   1127           4314         1468
## [3,]  27  9898   961   2861   3151           242          833
## [4,] 277 27901  3749   6964  4479           603         2503
## [5,] 251  3191  1993   1799   1730           234          710
## [6,] 302  5283 13316  20399   1809           8752          172
## [7,] 212 12119 28326  39694   4736          19410         2870
## [8,] 240 47493  2567   3779   5243           828         2253
## Clustering vector:
##   [1] 1 2 2 3 4 3 3 2 5 6 2 3 4 4 4 3 2 5 1 2 1 5 4 7 4 1 3 1 6 8 1 5 1 4 5
##  [36] 2 4 6 6 8 4 1 6 6 2 6 6 7 2 6 3 2 8 2 4 5 7 2 1 2 3 7 2 6 5 7 2 4 5 3
##  [71] 1 4 5 1 2 1 5 6 3 5 3 2 2 1 3 7 7 8 3 4 3 3 7 3 2 5 2 5 5 3 2 6 2 8 1
```

```
## [106] 1 2 6 2 6 3 6 1 1 1 3 3 3 1 3 1 5 3 2 4 8 1 4 5 8 3 5 1 3 3 5 2 2 1 5
## [141] 1 4 4 3 1 6 3 3 5 4 1 5 1 5 5 6 2 1 2 2 2 3 1 6 2 6 2 5 5 5 2 6 2 6 5
## [176] 2 8 3 3 5 3 8 2 7 5 3 5 2 2 2 1 1 5 2 3 1 4 2 3 3 6 6 4 5 5 6 5 2 5 6
## [211] 1 7 3 2 2 2 6 1 2 5 1 2 5 5 3 3 4 5 5 3 3 2 4 5 1 5 3 1 3 8 4 4 1 3 2
## [246] 2 3 3 1 3 5 7 3 4 2 4 3 3 8 8 3 3 1 5 2 6 6 1 6 1 5 5 2 4 5 5 4 3 3 1
## [281] 5 3 8 4 8 8 3 1 1 8 5 5 5 2 1 3 1 3 2 5 1 6 2 2 6 2 6 1 3 6 3 4 6 3 3
## [316] 2 3 5 3 6 5 3 1 1 4 4 5 5 1 5 3 6 1 7 1 4 1 3 5 5 2 2 2 6 5 2 2 4 5 6
## [351] 5 6 5 6 1 5 1 6 2 5 1 5 5 5 5 2 3 5 1 5 8 1 5 1 5 5 2 8 5 2 4 1 4 5 6
## [386] 3 5 1 3 3 5 5 5 4 3 3 2 3 3 3 5 4 4 4 1 3 4 6 3 3 3 5 2 3 3 2 2 2 6 3
## [421] 2 1 4 1 1 3 6 4 5 3 2 3 1 5 1 4 8 6 3 5
## Objective function:
##      build      swap
## 7099.316 6952.244
##
## Available components:
## [1] "medoids"      "id.med"      "clustering"  "objective"  "isolation"
## [6] "clusinfo"    "silinfo"    "diss"        "call"       "data"
```

```
plot(X,col=X.pam.8.clust$clustering)
```



```
plot(X.pam.8.clust, which.plots = 2)
```

Silhouette plot of pam(x = X, k = k, keep.diss = TRUE, keep.da

n = 440

8 clusters C_j j : n_j | $\text{ave}_{i \in C_j} s_i$

1 : 69 | 0.32

2 : 76 | 0.27

3 : 88 | 0.16


4 : 44 | 0.05

5 : 86 | 0.40

6 : 46 | 0.17

7 : 12 | -0.09


8 : 19 | 0.23



Average silhouette width : 0.24

```
# long lines good - means greater within cluster similarity
# Centroid Plot against 1st two discriminant functions
clusplot(X.pam.8.clust, color=TRUE, shade=TRUE, labels=2, lines=0)
```


clusplot(pam(x = X, k = k, keep.diss = TRUE, keep.data = TRUE))



These two components explain 72.46 % of the point variability.

1.2.4 Hierarchical Clustering

```
dfWCD2.h.clust<- hclust(d=dist(X))
plot(dfWCD2.h.clust, labels = FALSE)
rect.hclust(dfWCD2.h.clust, k=3, border="red")
rect.hclust(dfWCD2.h.clust, k=6, border="green")
rect.hclust(dfWCD2.h.clust, k=8, border="blue")
```

1.3 Questions:

1. How did you choose a model for EM? Evaluate the model performance. I used the same data from the previous assignment, Wholesale Customers Data Set. To format the data, I removed the text columns and the output columns. The output data is channel and region. This left the data Fresh, Milk, Grocery, Frozen, Detergents_Paper, Delicassen to be analyzed.

The cluster plot (clPairs) doesn't appear to cluster very well. The green is overlapping everything in all of the plots, it doesn't seem too separated from the rest of the data. The Mclust plot ended up clustering in 8 different tables instead of 3, like I expected from the region data. The classification table looks acceptable, from what I can tell, the cylinders (probability distributions) are showing the clusters well. In the density plots, it appears there is more variance in the estimates because the density plots are so large.

2. Cluster some of your data using EM based clustering that you also used for k-means, PAM, and hierarchical clustering. How do the clustering approaches compare on the same data?

When I set the clusters to 8 and re-run the M4L2 clustering, the results look similar. The k-means, pam, and EM cluster plots look identical. However, I do not think this data clusters well. Looking at the clPairs plot, the green is overlapping everything and when I look at the hierarchical clustering plot, a majority of the data falls into one cluster, even when set to 8 clusters.