

M6L1 Homework Assignment

Joshua Conte

October 29, 2017

1 M6L1 Homework Assignment

R studio was configured with the following parameters before beginning the project:

```
# clears the console in RStudio
cat("\014")
```

```
# clears environment
rm(list = ls())
```

```
# Load required packages
require(ggplot2)
require(class)
```

1.1 Load Data.

I opened the Wholesale customers Data Set using read.csv2 and downloaded it directly from the UC Irvine Machine Learning Repository.

To format the data, the data is separated by ',', stringsAsFactors = FALSE so that the strings in a data frame will be treated as plain strings and not as factor variables. I set na strings for missing data. Once the data was loaded I added the column names and changed the data types to numeric and finally removed the text data type.

Below is my R code:

```
# Some csv files are really big and take a while to open. This command checks to
# see if it is already opened, if it is, it does not open it again.
# I also omitted the first column
if (!exists("dfWCD")) {
dfWCD <-
  read.csv2("Wholesale customers data.csv",
    sep = ",",
    stringsAsFactors = FALSE,
    na.strings=c("", "NA")
  )
}

# Download directly from site (unreliable from Ecuador)
# if (!exists("dfWCD")) {
# dfWCD <-
#   read.csv2(
#     url(
#       "https://archive.ics.uci.edu/ml/machine-learning-databases/00292/Wholesale customers data.csv"
#     ),
#     sep = ",",
#     stringsAsFactors = FALSE,
#     na.strings=c("", "NA")
#   )
# # Add a column so I know which study the data is referring to
# study <- sprintf("study_%s", seq(1:440))
# dfWCD$study<-study
# }

# change 2 to 24 to numeric
```

```
dfWCD[1:8] <- sapply(dfWCD[1:8], as.numeric)

# Print first lines
str(dfWCD)

## 'data.frame': 440 obs. of 8 variables:
## $ Channel      : num  2 2 2 1 2 2 2 1 2 ...
## $ Region       : num  3 3 3 3 3 3 3 3 3 ...
## $ Fresh        : num  12669 7057 6353 13265 22615 ...
## $ Milk         : num  9656 9810 8808 1196 5410 ...
## $ Grocery      : num  7561 9568 7684 4221 7198 ...
## $ Frozen       : num  214 1762 2405 6404 3915 ...
## $ Detergents_Paper: num  2674 3293 3516 507 1777 ...
## $ Delicassen   : num  1338 1776 7844 1788 5185 ...
```

1.1.1 Understanding the data

The data set refers to clients of a wholesale distributor in Portugal. It includes the annual spending in monetary units (m.u.) on diverse product categories. The data has the following attribute information:

1. FRESH: annual spending (m.u.) on fresh products (Continuous);
2. MILK: annual spending (m.u.) on milk products (Continuous);
3. GROCERY: annual spending (m.u.) on grocery products (Continuous);
4. FROZEN: annual spending (m.u.) on frozen products (Continuous)
5. DETERGENTS_PAPER: annual spending (m.u.) on detergents and paper products (Continuous)
6. DELICATESSEN: annual spending (m.u.) on and delicatessen products (Continuous);
7. CHANNEL: customer channel - 1 = Horeca (Hotel/Restaurant/Cafe) or 2 = Retail
8. REGION: Customers Region - 1= Lisbon 2 = Oporto or 3 = Other (Nominal)

1.2 k-Nearest Neighbors (kNN) in R

A simple supervised learning algorithm is k-Nearest Neighbors algorithm (k-NN). KNN is a non-parametric method used for classification and regression.

In both cases, the input consists of the k closest training examples in the feature space. The output depends on whether k-NN is used for classification or regression:

In k-NN classification, the output is a class membership. An object is classified by a majority vote of its neighbors, with the object being assigned to the class most common among its k nearest neighbors (k is a positive integer, typically small). If $k = 1$, then the object is simply assigned to the class of that single nearest neighbor.

I am using the Channel data for classification:

```
head(dfWCD)

##   Channel Region Fresh Milk Grocery Frozen Detergents_Paper Delicassen
## 1       2      3 12669 9656   7561    214           2674         1338
## 2       2      3  7057 9810   9568   1762           3293         1776
## 3       2      3  6353 8808   7684   2405           3516         7844
## 4       1      3 13265 1196   4221   6404            507         1788
## 5       2      3 22615 5410   7198   3915           1777         5185
## 6       2      3  9413 8259   5126    666           1795         1451
```

```
summary(dfWCD)
```

```
##      Channel      Region      Fresh      Milk
##  Min.   :1.000   Min.   :1.000   Min.    :    3   Min.    :   55
## 1st Qu.:1.000   1st Qu.:2.000   1st Qu.: 3128   1st Qu.: 1533
##  Median:1.000   Median:3.000   Median:  8504   Median: 3627
##  Mean   :1.323   Mean   :2.543   Mean    :12000   Mean    : 5796
## 3rd Qu.:2.000   3rd Qu.:3.000   3rd Qu.:16934   3rd Qu.: 7190
##  Max.   :2.000   Max.    :3.000   Max.    :112151   Max.    :73498
##  Grocery      Frozen      Detergents_Paper  Delicassen
##  Min.    :    3   Min.    :   25.0   Min.    :    3.0   Min.    :    3.0
## 1st Qu.: 2153   1st Qu.: 742.2   1st Qu.: 256.8   1st Qu.: 408.2
##  Median: 4756   Median: 1526.0   Median:  816.5   Median:  965.5
##  Mean    : 7951   Mean    :3071.9   Mean    :2881.5   Mean    :1524.9
## 3rd Qu.:10656   3rd Qu.:3554.2   3rd Qu.:3922.0   3rd Qu.:1820.2
##  Max.    :92780   Max.    :60869.0   Max.    :40827.0   Max.    :47943.0
```

```
length(dfWCD)
```

```
## [1] 8
```

```
names(dfWCD)
```

```
## [1] "Channel"      "Region"      "Fresh"
## [4] "Milk"         "Grocery"     "Frozen"
## [7] "Detergents_Paper" "Delicassen"
```

```
table(dfWCD$Channel)
```

```
##
##  1  2
## 298 142
```

```
dfWCD$Channel
```

```
## [1] 2 2 2 1 2 2 2 2 1 2 2 2 2 2 2 2 1 2 1 2 1 2 1 1 2 2 2 1 1 2 1 1 1 1 1
## [36] 2 1 2 2 1 1 1 2 2 2 2 2 2 2 2 1 1 2 2 1 1 2 2 1 1 2 2 2 2 1 2 1 2 1 1
## [71] 1 1 1 2 2 1 1 2 1 1 1 2 2 1 2 2 2 1 1 1 1 2 1 2 1 2 1 1 1 2 2 2 1 1
## [106] 1 2 2 2 2 1 2 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1
## [141] 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 2 2 1 2 2 2 1 1 2 2 2 2 1 1 1 2 2 1 2 1
## [176] 2 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 1 1 1 2 1 1 1 2 1 1 2 2 1 1 1 2 1 2 1 2
## [211] 1 2 1 1 2 1 2 1 2 1 1 1 1 2 1 1 2 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [246] 2 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 2 1 2 1 2 1 1 1 1 1 1 1 1 1 1 2
## [281] 1 2 1 1 1 1 1 1 1 1 1 1 1 2 1 2 1 2 2 1 2 2 2 2 2 2 2 2 1 1 2 1 1 2 1 1
## [316] 2 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 2 1 2 2 2 1 1 1 1 2 2 1 2 1 1 2 2 1 2
## [351] 1 2 1 2 1 1 1 2 1 1 1 1 1 1 1 2 1 1 1 1 2 1 1 2 1 1 2 1 1 2 1 1 1 1
## [386] 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 2 2 1 1 1 1 1 1 2 2 1 2 1
## [421] 1 2 1 2 2 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1
```

```
length(dfWCD$Channel)
```

```
## [1] 440
```

In this section I shuffle the data to make it a little more random to make sure the training set and predictive set have a good random sample, I don't want to train on all 1's to predict for the 2's:

```
shuff<-runif(nrow(dfWCD))
head(shuff)
```


```
## [1] 0.05098633 0.28433827 0.36240784 0.79634407 0.46410232 0.86816866
```

```
shopping<-dfWCD[order(shuff),]
shopping$Channel
```


```
## [1] 1 2 2 2 1 1 1 1 1 1 2 1 1 2 2 1 1 2 2 1 1 2 1 1 1 2 1 1 1
## [36] 1 2 1 2 1 1 2 2 1 2 1 1 2 2 1 2 1 1 1 1 1 1 2 2 1 2 1 1 1 1 2
## [71] 1 1 1 2 1 2 2 2 2 1 1 1 2 1 1 2 1 2 1 1 1 2 1 1 1 2 2 2 1 2 1 1 1
## [106] 1 2 2 1 1 1 2 1 1 1 1 2 2 1 1 1 2 1 1 1 1 1 1 2 1 2 2 2 1 2 1 2 1 1
## [141] 1 2 1 2 1 1 2 1 1 1 1 2 1 1 1 2 1 1 2 1 1 1 1 1 1 2 1 1 1 2 1 1 2 1
## [176] 2 2 1 2 1 2 1 2 1 1 2 1 1 2 1 1 1 1 2 1 1 1 1 2 1 2 2 1 1 1 1 2 2
## [211] 2 1 1 2 1 2 1 1 1 2 2 1 2 1 2 1 2 1 1 2 1 2 1 2 1 1 1 1 2 1 1 1 1
## [246] 1 2 1 1 2 1 1 1 2 2 1 2 1 2 1 1 1 1 2 1 1 1 2 1 1 1 1 2 1 1 2 1 1 1
## [281] 2 1 1 1 2 1 1 2 1 1 1 1 2 1 2 1 2 1 2 2 1 1 2 1 1 1 2 1 1 1 1 2 2 1
## [316] 2 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 2 1 1 2 1 2 1 1 1
## [351] 1 1 1 1 1 2 1 1 2 1 1 1 2 1 2 1 1 2 1 2 1 2 1 1 2 1 2 1 1 1 1 2 1 1
## [386] 2 2 2 1 1 1 1 1 2 1 1 2 1 1 1 2 1 2 1 2 2 1 1 1 2 1 1 1 2 1 1 1 1
## [421] 1 1 2 1 2 2 1 1 2 1 1 2 1 1 1 1 1 1 1 2
```

This looks at the plotting to see if the knn will do a good job. Here I plot Detergents_Paper vs Milk and Detergents_Paper vs Fresh:

```
qplot(shopping$Detergents_Paper,shopping$Milk,data=shopping)+geom_point(aes(colour = factor(shopping$Channel
```



```
qplot(shopping$Detergents_Paper,shopping$Fresh,data=shopping)+geom_point(aes(colour = factor(shopping$Ch
```



```
summary(shopping)
```

```
##      Channel      Region      Fresh      Milk
##  Min.   :1.000   Min.   :1.000   Min.    :    3   Min.    :   55
## 1st Qu.:1.000   1st Qu.:2.000   1st Qu.: 3128   1st Qu.: 1533
## Median :1.000   Median :3.000   Median : 8504   Median : 3627
## Mean   :1.323   Mean   :2.543   Mean   :12000   Mean   : 5796
## 3rd Qu.:2.000   3rd Qu.:3.000   3rd Qu.:16934   3rd Qu.: 7190
## Max.   :2.000   Max.   :3.000   Max.   :112151   Max.   :73498
##      Grocery      Frozen      Detergents_Paper      Delicassen
##  Min.    :    3   Min.    :   25.0   Min.    :    3.0   Min.    :    3.0
## 1st Qu.: 2153   1st Qu.:  742.2   1st Qu.:  256.8   1st Qu.:  408.2
## Median : 4756   Median : 1526.0   Median :   816.5   Median :   965.5
## Mean    : 7951   Mean    : 3071.9   Mean    :  2881.5   Mean    : 1524.9
## 3rd Qu.:10656   3rd Qu.: 3554.2   3rd Qu.: 3922.0   3rd Qu.: 1820.2
## Max.    :92780   Max.    :60869.0   Max.    :40827.0   Max.    :47943.0
```

This algorithm should do a good job since the data is not mixed, the plots look separated.

Here I grab the data that can be analyzed, I no longer need Channel or Regions, and I scale the data to make it all on a comparable scale.

```
shopping.scaled<-as.data.frame(lapply(shopping[,c(3:8)], scale))
head(shopping.scaled)
```

```
##      Fresh      Milk      Grocery      Frozen Detergents_Paper
## 1 -0.054264243 -0.3666840 -0.6197176  6.57862352   -0.5894671
## 2 -0.747928502 -0.1435246 -0.1959639 -0.02594033    0.5106085
```

```
## 3 -0.591848113 0.8393520 0.1466588 -0.08814843 0.1041363
## 4 -0.007297804 0.3057749 1.2517646 -0.41257808 0.7325112
## 5 -0.458460264 0.2681074 0.3015546 -0.38909556 0.0548479
## 6 1.022247655 -0.5413363 -0.5762584 -0.17631090 -0.5179464
## Delicassen
## 1 0.4159878
## 2 -0.4981623
## 3 0.2440084
## 4 -0.2932055
## 5 0.1578414
## 6 -0.1694512
```

```
summary(shopping.scaled)
```

```
##      Fresh      Milk      Grocery      Frozen
## Min.   :-0.9486 Min.   :-0.7779 Min.   :-0.8364 Min.   :-0.62763
## 1st Qu.: -0.7015 1st Qu.: -0.5776 1st Qu.: -0.6101 1st Qu.: -0.47988
## Median :-0.2764 Median :-0.2939 Median :-0.3363 Median :-0.31844
## Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.00000
## 3rd Qu.: 0.3901 3rd Qu.: 0.1889 3rd Qu.: 0.2846 3rd Qu.: 0.09935
## Max.    : 7.9187 Max.    : 9.1732 Max.    : 8.9264 Max.    :11.90545
## Detergents_Paper Delicassen
## Min.   :-0.6037 Min.   :-0.5396
## 1st Qu.: -0.5505 1st Qu.: -0.3960
## Median :-0.4331 Median :-0.1984
## Mean   : 0.0000 Mean   : 0.0000
## 3rd Qu.: 0.2182 3rd Qu.: 0.1047
## Max.    : 7.9586 Max.    :16.4597
```

This also scales the data like above and puts all of it on a scale from 0 to 1.

```
normalize<- function(x) {
  return((x-min(x))/(max(x)-min(x)))
}
shopping.normalized<-as.data.frame(lapply(shopping[,c(3:8)],normalize))
head(shopping.normalized)
```

```
##      Fresh      Milk      Grocery      Frozen Detergents_Paper Delicassen
## 1 0.10085780 0.04132456 0.02219300 0.57497863 0.001665687 0.056216103
## 2 0.02263081 0.06375012 0.06559815 0.04800802 0.130144033 0.002440551
## 3 0.04023255 0.16252059 0.10069306 0.04304451 0.082671958 0.046099291
## 4 0.10615437 0.10890078 0.21388922 0.01715864 0.156060161 0.014497288
## 5 0.05527517 0.10511553 0.11655906 0.01903228 0.076915540 0.041030455
## 6 0.22225987 0.02377354 0.02664453 0.03601012 0.010018617 0.021777222
```

```
summary(shopping.normalized)
```

```
##      Fresh      Milk      Grocery      Frozen
## Min.   :0.00000 Min.   :0.00000 Min.   :0.00000 Min.   :0.00000
## 1st Qu.:0.02786 1st Qu.:0.02012 1st Qu.:0.02317 1st Qu.:0.01179
## Median :0.07580 Median :0.04864 Median :0.05122 Median :0.02467
## Mean   :0.10698 Mean   :0.07817 Mean   :0.08567 Mean   :0.05008
## 3rd Qu.:0.15097 3rd Qu.:0.09715 3rd Qu.:0.11482 3rd Qu.:0.05800
## Max.    :1.00000 Max.    :1.00000 Max.    :1.00000 Max.    :1.00000
## Detergents_Paper Delicassen
## Min.   :0.000000 Min.   :0.000000
## 1st Qu.:0.006216 1st Qu.:0.008453
```

```
## Median :0.019927 Median :0.020077
## Mean   :0.070510 Mean    :0.031745
## 3rd Qu.:0.095997 3rd Qu.:0.037907
## Max.    :1.000000 Max.    :1.000000
```

```
nrow(shopping)
```

```
## [1] 440
```

This takes the normalized data and trains the first 390 to predict the remaining 50, with a k=4:

```
shopping.normalized.train<-shopping.normalized[1:390,]
shopping.normalized.test<-shopping.normalized[391:440,]
shopping.normalized.train.target<-shopping[1:390,c(1)]
shopping.normalized.test.target<-shopping[391:440,c(1)]
shopping.normalized.test.target
```

```
## [1] 1 1 1 2 1 1 2 1 1 1 2 1 2 1 2 2 1 1 1 2 1 1 1 2 1 1 1 2 1 2
## [36] 2 1 1 2 1 1 2 1 1 1 1 1 1 1 2
```

```
k<-4
```

```
knn.m1.4<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train,
knn.m1.4
```

```
## [1] 1 1 1 2 1 2 2 1 1 1 2 1 2 1 2 2 1 1 1 2 1 1 1 2 1 1 1 2 1 2 2 2 2
## [36] 2 1 1 2 1 1 2 1 1 1 1 1 1 1 2
## Levels: 1 2
```

```
length(knn.m1.4)
```

```
## [1] 50
```

```
cm<-table(shopping.normalized.test.target,knn.m1.4)
cm
```

```
##                knn.m1.4
## shopping.normalized.test.target  1  2
##                                1 31  4
##                                2  0 15
```

The results look like 3 out of 33 were mislabeled as Channel 2 instead of Channel 1 (9.09% error) and 0 out of 17 were mislabeled as Channel 1 instead of Channel 2 (0% error). Overall, I think it did a good job.

1.2.1 Using different k's

This uses k=2:

```
k<-2
knn.m1.2<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train,
cm<-table(shopping.normalized.test.target,knn.m1.2)
cm
```

```
##                knn.m1.2
## shopping.normalized.test.target  1  2
##                                1 32  3
##                                2  2 13
```

Using k = 2 gives a great result of 1 out of 33 were mislabeled as Channel 2 instead of Channel 1 (3.03% error) and 0 out of 17 were mislabeled as Channel 1 instead of Channel 2 (0% error).

This uses $k=3$:

```
k<-3
knn.m1.3<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train, k=3)
cm<-table(shopping.normalized.test.target,knn.m1.3)
cm
```

```
##                knn.m1.3
## shopping.normalized.test.target  1  2
##                1 32  3
##                2  0 15
```

Using $k = 3$ gives a great result of 2 out of 33 were mislabeled as Channel 2 instead of Channel 1 (6.06% error) and 0 out of 17 were mislabeled as Channel 1 instead of Channel 2 (0% error).

This uses $k=5$:

```
k<-5
knn.m1.5<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train, k=5)
cm<-table(shopping.normalized.test.target,knn.m1.5)
cm
```

```
##                knn.m1.5
## shopping.normalized.test.target  1  2
##                1 30  5
##                2  1 14
```

Using $k = 5$ gives the same result as $k = 4$.

This uses $k=6$:

```
k<-6
knn.m1.6<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train, k=6)
cm<-table(shopping.normalized.test.target,knn.m1.6)
cm
```

```
##                knn.m1.6
## shopping.normalized.test.target  1  2
##                1 30  5
##                2  0 15
```

Using $k = 6$ gives the same result as $k = 4$.

This uses $k=7$:

```
k<-7
knn.m1.7<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train, k=7)
cm<-table(shopping.normalized.test.target,knn.m1.7)
cm
```

```
##                knn.m1.7
## shopping.normalized.test.target  1  2
##                1 30  5
##                2  1 14
```

Using $k = 7$ gives the same result as $k = 3$.

This uses $k=8$:

```
k<-8
knn.m1.8<-knn(train = shopping.normalized.train, test = shopping.normalized.test,shopping.normalized.train, k=8)
```

```
cm<-table(shopping.normalized.test.target,knn.m1.8)
cm
```

```
##                      knn.m1.8
## shopping.normalized.test.target  1  2
##                      1 31  4
##                      2  0 15
```

Using $k = 8$ gives the same result as $k = 3$.

1.2.2 Using the original data

Does scaling, normalization or leaving the data unscaled make a difference for kNN?

```
shopping.train<-shopping[1:390,]
shopping.test<-shopping[391:440,]
shopping.train.target<-shopping[1:390,c(1)]
shopping.test.target<-shopping[391:440,c(1)]
shopping.test.target
```

```
## [1] 1 1 1 2 1 1 2 1 1 1 2 1 2 1 2 2 1 1 1 2 1 1 1 2 1 1 1 1 1 2 1 2
## [36] 2 1 1 2 1 1 2 1 1 1 1 1 1 1 1 2
```

```
k<-3
knn.m1.4<-knn(train = shopping.train, test = shopping.test,shopping.train.target,k)
knn.m1.4
```

```
## [1] 1 2 1 2 1 2 2 1 1 1 2 1 2 1 2 2 1 1 1 2 1 1 1 2 1 1 2 2 1 2 2 2 2
## [36] 2 1 1 2 1 1 2 1 1 1 1 1 1 1 1 2
## Levels: 1 2
```

```
length(knn.m1.4)
```

```
## [1] 50
```

```
cm<-table(shopping.test.target,knn.m1.4)
cm
```

```
##                      knn.m1.4
## shopping.test.target  1  2
##                      1 29  6
##                      2  0 15
```

No, because the data was already to scale.

1.3 Questions

- Does the k for kNN make a difference? Try for a range of values of k .
 - Yes, the k value makes a difference. I ran KNN with a k value from 2 through 8, 2 was the best followed by 3. A k value of 4, 5, and 6 gave the worst and a k value of 7 and 8 were the same of a k value of 3.
- Does scaling, normalization or leaving the data unscaled make a difference for kNN? Why or Why not?
 - For this dataset it does not matter because all of the data used were at a similar scale. If I used a different dataset that had mixed data at different scales, like the wine data used in the class lecture, it would make a difference. This is because distance metrics that are not scale invariant, attributes with larger absolute ranges will dominate.